

## Problem A. NJU Emulator

First, we consider the idea of doubling, for any  $N \geq 2$ , we can obtain  $N$  using  $(\times 2)$  or  $(\times 2 + 1)$  from 1 in  $\lfloor \log_2 N \rfloor$  steps. Then we may come up with an algorithm to construct any  $N$  within  $O(\log N)$  s86 instructions as follows:

Use 3 instructions "p1, dup, add 1" to push 1 and 2 into stack, then use "p1" and  $(\times 2)$  or  $(\times 2 + 1)$  for  $\lfloor \log_2 N \rfloor$  times to get  $N$ .

It takes at most  $2\lfloor \log_2 N \rfloor + 5$  instructions to construct  $N$ , which is too large. To improve our algorithm, we can consider applying the method of four Russians, i.e., try to use a larger base  $k$  instead of the base 2.

For any  $k \geq 2$  and  $N \geq 2$ , we can obtain  $N$  using  $(\times k + p)$  ( $0 \leq p < k$ ) from some  $x$  ( $1 \leq x < k$ ) by  $\lfloor \log_k N \rfloor$  steps. Thus we can use  $2k - 1$  instructions "p1, dup, add 1, dup, add 1, ..." to push  $1, 2, \dots, k$  into the stack, then use "p1, add (x-1)" and  $\lfloor \log_k N \rfloor$  times of  $(\times k + p)$  ( $0 \leq p < k$ ) operation to get  $N$ .

It takes at most  $2\lfloor \log_k N \rfloor + 2k + 2$  instructions to construct  $N$ . Take  $k = 7$  (or 8, 9, 10, 11, 12), we obtain an algorithm to construct any  $N < 2^{64}$  within 60 s86 instructions, which is close to our goal.

For further improvements, note that for a base  $k$ , if we have obtained  $1, 2, \dots, \lfloor k/2 \rfloor$  and  $k$  in the stack, then for any  $x \leq k$ , we can get  $x$  in 2 steps as following:

- "p1, add  $x - 1$ " when  $x \leq \lfloor k/2 \rfloor$ .
- "dup, sub  $k - x$ " when  $x > \lfloor k/2 \rfloor$ .

Then for  $N > k$ , we can obtain  $N$  using  $(\times k + p)$  ( $0 \leq p \leq \lfloor k/2 \rfloor$ ) or  $(\times k - p)$  ( $1 \leq p \leq \lfloor k/2 \rfloor$ ) from some  $x$  ( $1 \leq x \leq k$ ) in  $\lfloor \log_k N \rfloor$  steps. Thus we can use  $k + 1$  instructions to push  $1, 2, \dots, \lfloor k/2 \rfloor$  and  $k$  into the stack, then construct  $x$  in 2 steps, finally applying  $\lfloor \log_k N \rfloor$  times of  $(\times k + p)$  ( $0 \leq p \leq \lfloor k/2 \rfloor$ ) or  $(\times k - p)$  ( $1 \leq p \leq \lfloor k/2 \rfloor$ ) to get  $N$ .

It takes at most  $2\lfloor \log_k N \rfloor + k + 4$  instructions to construct  $N$ . Take  $k = 16$  (or 12, 14), we can obtain an algorithm to construct any  $N < 2^{64}$  within 50 s86 instructions.

## Problem B. Just another board game

Consider what is the optimal strategy if the "instant-ending" operation is not allowed. We can reinterpret the operations of the two players as follows:

We start with  $r = 1$  and  $c = 1$ . In each turn, if it's Roundgod's turn to move, he can set  $c$  to some arbitrary value satisfying  $1 \leq c \leq m$ . If it's kimoyami's turn to move, he can set  $r$  to some arbitrary value satisfying  $1 \leq r \leq n$ . After  $k$  rounds, the game ends, and the value of the game is  $a_{r,c}$ .

We consider the process backward. If it's Roundgod to move in the last turn, he should obviously choose  $c$  to be the position of the maximum element in the current row, and thus kimoyami, who is second-to-last to move, should choose  $r$  to be the row with the minimal maximum. Similarly, if it's kimoyami to move in the last turn, he should choose  $r$  to be the position of the minimum element in the current column, and thus Roundgod, who is second-to-last to move, should choose  $c$  to be the row with the maximal minimum.

There still exists a special case where  $k = 1$ , i.e., the game only lasts for at most one turn, then clearly it's optimal for Roundgod to choose the maximum in the first row.

So, let's now consider the case where the "instant-ending" operation is allowed. I claim that this operation wouldn't make much difference to the strategy for both players, except for the first move of Roundgod.

Why is this the case? Because if you choose the second operation after your opponent has moved, it is never better than doing nothing (keeping the chess unmoved), as in the optimal strategy of your opponent, he would also do nothing (or if he chooses the second operation, it's still the same), otherwise, if he chooses

another  $r$  (or  $c$ ) different than his last move, which is clearly not optimal, and thus you can still have the same value without doing the second operation.

To summarize, the answer can be calculated as follows:

Let  $rowmax_i$  be the maximum element in the  $i$ th row, and  $colmin_i$  be the minimum element in the  $i$ th column.

- If  $k = 1$ , then the answer is  $rowmax_1$
- Otherwise, if  $k$  is odd, then the answer is  $\max(a_{1,1}, \min_{i=1}^n rowmax_i)$
- Otherwise,  $k$  is even, and the answer is  $\max(a_{1,1}, \max_{i=1}^m colmin_i)$

The overall time complexity is  $O(nm)$ .

## Problem C. Dota2 Pro Circuit

First, we show how to compute  $worst_i$  for each team, and  $best_i$  for each team can be computed in a similar manner.

For fixed  $i$  and  $k$ , in order to check whether there exists a plan that  $k$  teams have a higher score than team  $i$ , we will let team  $i$  scores  $b_n$  in the tournament, and the  $k$  teams with the initial highest score (exclude team  $i$ ) getting  $b_1, b_2, \dots, b_k$  scores in the tournament. We will match the initial scores and tournament scores in a greedy way: the team with the highest initial score gets  $b_k$  scores in the tournament and the team with the second-highest initial score get  $b_{k-1}$  scores in the tournament, and so on. Then we check if all selected  $k$  teams' final scores are strictly higher than team  $i$ . This takes  $O(n)$  time.

In order to compute the answer for all  $i$ , using binary search would take  $O(n^2 \log n)$  time, which would be too slow. An obvious observation is that, if we sort the teams by their initial scores, then  $worst_i$  is non-increasing. Thus we can use the method of two pointers to achieve a total complexity of  $O(n^2)$ .

## Problem D. Into the Woods

For a point  $P(x, y)$ , the Manhattan distance between  $P$  and  $O(0, 0)$  is  $|x| + |y| = \max\{|x + y|, |x - y|\}$ . Assume Roundgod walks through a path  $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$ . Let  $A$  be the maximum Manhattan distance between any point of the path to  $(0, 0)$ , then  $A = \max\{\max_{i=1}^n |x_i + y_i|, \max_{i=1}^n |x_i - y_i|\}$ .

By the linearity of expectation, what we need to calculate is

$$E(A) = \sum_{k=1}^n Pr[A \geq k] = n - \sum_{k=1}^n Pr[A \leq k]$$


Here's a very useful (and beautiful) trick for this kind of two-dimensional random walk: Define random variables  $a = x + y$  and  $b = x - y$ , then  $a$  and  $b$  are independent random variables with the same distribution! Formally, assume Roundgod's position is  $(x, y)$ . In each step,  $x + y$  and  $x - y$  will independently increase or decrease by 1 with  $\frac{1}{2}$  probability.

Now the problem is reduced as follows: We go on a one-dimensional random walk on the  $x$ -axis for  $n$  steps, starting from  $x = 0$ , each step randomly setting  $x = x + 1$  or  $x = x - 1$  with probability  $\frac{1}{2}$  each. Let  $p_k$  denote the probability that the path of the random walk is completely inside interval  $[-k, k]$ . Compute  $p_k$  for each  $k = 1, 2, \dots, n - 1$ . Then the answer is

$$E(A) = n - \sum_{k=1}^{n-1} p_k^2$$

We can now transform this to an equivalent problem: suppose one starts from  $(0, 0)$ , walks for  $n$  steps, with each step only allowed to go right or go up. Let  $c_k$  be the number of paths which don't collide with line  $y = x + (k + 1)$  or the line  $y = x - (k + 1)$ , then  $p_k = \frac{c_k}{2^n}$ . To compute this number, one still need to solve another subproblem: Given a point  $P$ , line  $l_0 : y = x + a$  and line  $l_1 : y = x - b$ , how many paths from  $(0, 0)$  to  $P$  are there that don't trespass line  $l_0$  or line  $l_1$ ? One may recall a similar but simpler problem(also known as a generalization Bertrand's ballot theorem): Given a point  $P$  and a line  $l : y = x + a$ , how many paths from  $(0, 0)$  to  $P$  are there that doesn't trespass line  $l$ ? Note that if  $a = 0$ , then this number is known to be the Catalan number.

So now we finally reveal the key idea to this problem, which I think is strikingly beautiful. It's called André's reflection principle, whose main idea is captured by the following picture.



By the reflection principle, we can see that the answer when only one restricting line exists is the difference between two binomial coefficients. And by an inclusive and exclusive version of reflection principle, let  $f(P)$  be the mirror reflection point of  $P$  and line  $l_0$ ,  $g(P)$  be the mirror reflection point of  $P$  and line  $l_1$ ,  $C(P)$  be the number of ways of going to point  $P$  starting from  $(0, 0)$ , then the answer is

$$C(P) - C(f(P)) - C(g(P)) + C(g(f(P))) + C(f(g(P))) - \dots$$

.

Thus we can compute  $c_k$  as

$$c_k = \sum_{i=L}^R \binom{n}{i} + 2 \sum_{j \geq 1} (-1)^j \sum_{i=L}^R \binom{n}{i - (k+1)j},$$

where  $L = \lceil \frac{n-k}{2} \rceil$  and  $R = \lfloor \frac{n-k}{2} \rfloor$ .

If we preprocess the prefix sum of  $\binom{n}{i}$ , then  $c_k$  can be compute in  $O(\frac{n}{k})$  time, and thus the overall time complexity is  $\sum_{i=1}^n \frac{n}{i} = O(n \log n)$ .

## Problem E. Did I miss the lethal?

This problem can be reinterpreted as a game:

There are  $n$  cards in your hand, and the  $i$ th card has corresponding  $d_i$  and  $a_i$  on it. You and an opponent(the dealer, that's what we call it in Hearthstone) take turns to move in each round:

1. You choose a card with  $d_i$  and  $a_i$  on it, remove the card and deal  $d_i$  damage.
2. The opponent chooses  $a_i$  cards and removes them.

You want to maximize the total damages you deal while the opponent wants to minimize it, you should compute the answer if you and the opponent play optimally.

Note that each  $a_i$  is in  $\{1, 2, 3, 4\}$ , thus we can divide all  $n$  cards into 4 piles: all cards with  $d_i = j$  is split into the  $j$ th pile ( $1 \leq j \leq 4$ ). Let  $n_1, n_2, n_3, n_4$  denote the number of cards in each pile initially.

Clearly, a greedy strategy follows: Once you or the opponent decide to remove a card from a pile, the card with the highest damage in that pile is chosen. Therefore for each pile, we can sort all cards by their  $d_i$ , and we can solve the problem using a dynamic programming approach:

Let  $dp[x_1][x_2][x_3][x_4][0]$  denote the answer where the  $j$ th pile remains  $x_j$  cards for each  $1 \leq j \leq 4$  and it is your turn to choose a card. Let  $dp[x_1][x_2][x_3][x_4][k]$  ( $k = 1, 2, 3, 4$ ) denote the answer where the  $j$ th pile remains  $x_j$  cards for each  $1 \leq j \leq 4$  and it is the opponent's turn to choose  $k$  cards. The state transitions can be computed in  $O(1)$ .

The total number of states is  $5n_1n_2n_3n_4$ , where  $\sum_{i=1}^4 n_i \leq 200$ , Thus the number of states reaches the maximum when  $n_1 = n_2 = n_3 = n_4 = 50$ , which is  $50^4 * 5$ , and should easily fit into time limit.

## Problem F. Guess The Weight

First, consider the optimal strategy for `uuzlovotree`, which is easy: after drawing the first card, he should look at the remaining cards in the deck and choose "Small" or "Large" based on which side has a greater number of cards compared to the first card he drew is.

A crucial observation is that, for any deck, there exists some "midpoint"  $mp$ , such that in the optimal strategy, `uuzlovotree` should choose "Large" if the first card he drew has a number less than or equal to  $mp$ , and should choose "Small" if the first card he drew has a number greater than  $mp$ .

Assume that the current multiset of cards is  $SC$ , and the number of cards in the deck is  $s$  (i.e.,  $|SC| = s$ ). Also we define  $L = \{x | x \in SC \wedge x \leq mp\}$ ,  $R = \{x | x \in SC \wedge x > mp\}$  then we can write the probability of drawing the second card as

$$Pr[DrawSecond] = 1 - \frac{\sum_{x \in L, y \in L} [x \geq y] + \sum_{x \in R, y \in R} [x \leq y]}{s(s+1)},$$

by considering all possible combination of the first two cards on the top of the deck.

Thus, by maintaining the number of cards and the number of pairs of cards with different values in each interval in a segment tree, we can calculate this probability in  $O(\log \max a_i)$ . Also, the midpoint can be found in this time using a descent on the segment tree. The overall time complexity is  $O(n + \max a_i + q \log \max a_i)$ .

## Problem G. Boring data structure problem

Consider maintaining the left side (left to the element in the middle) and the right side (right to the element in the middle) by two doubly-linked lists  $L$  and  $R$ , so that the first two insertion operations can both be done in  $O(1)$  time per query.

After an operation has been done, adjust the sizes of the two doubly linked lists as follows:

- If  $L > R$ , move the last element of  $L$  to the beginning of  $R$
- If  $L < R + 1$ , move the first element of  $R$  to the end of  $L$

This adjustment clearly has  $O(1)$  amortized complexity per query.

When deleting the element, we only need to mark the element as deleted and decrease the size of the corresponding doubly linked list by one. This is known as a 'lazy' operation, we can propagate the deletion

only when we are asked to perform operations(move or query) on it, which makes deletion also amortized  $O(1)$  complexity.

Therefore, the overall time complexity is  $O(n + q)$ .

## Problem H. Integers Have Friends 2.0

First, observe that if we take  $m = 2$  and either taking all odd  $a_i$  or all even  $a_i$ , we are guaranteed to have an answer at least  $\lceil \frac{n}{2} \rceil$ .

Randomly choose two distinct indices  $1 \leq x < y \leq n$ , the probability that the optimal answer contains both of them is at least  $\frac{1}{4}$ . If they are both contained in the answer, then  $m$  must be a factor of  $|a_x - a_y|$ . In fact, we only need to check those prime factors of  $|a_x - a_y|$ . Every positive  $x$  satisfying  $x \leq 4 \times 10^{12}$  has at most 11 distinct prime factors, and finding all of them can be easily done in  $O(\sqrt{\max a_i})$  time if precomputing all prime numbers up to  $2 \times 10^6$ .

Repeat the process above for  $K$  times. We will err with probability at most  $(\frac{3}{4})^K$ . The running time will be  $O(K\sqrt{\max a_i} + 11Kn)$ , taking  $K = 40$  should both has a negligible probability of error and easily fits into the time limit.

## Problem I. Little Prince and the garden of roses

Clearly, we can solve the problem for each color of the roses independently. Fix a certain color  $c$ . We need to assign colors to (the stems) of roses with color  $c$ . A common reduction in the grid works as follows: Consider the  $n \times n$  grid as a bipartite graph  $G = (V = L \cup R, E)$  with  $|L| = |R| = n$ , and consider each rose at position  $(u, v)$  as an edge between the  $u$ th vertex in the left and the  $v$ th vertex in the right part. One can find out this actually becomes the problem of edge coloring of bipartite graphs(i.e., give each edge a color, such that no two edges incident to the same vertex have the same color, minimize the number of distinct colors used).

So how many colors do we need? Let  $\Delta$  be the maximum degree of a vertex in the bipartite graph, then obviously we need at least  $\Delta$  colors. We can show that  $\Delta$  colors are also sufficient for bipartite graphs(Vizing's Theorem: The edge chromatic number of a general graph is  $\Delta$  or  $\Delta + 1$ ). There are two possible approaches:

First, we show a constructive algorithm that works in  $O(|E||V|)$ . Let's color the edges one by one in some order. Let  $(x, y)$  be the current edge. If there exists color  $c$  that is free in vertex  $x$  and vertex  $y$  then we can simply color  $(x, y)$  with  $c$ . If there is no such color then there are a couple of colors  $c_1, c_2$  so that  $c_1$  is in  $x$  and not in  $y$ , while  $c_2$  is in  $y$  but not in  $x$ . Let's make vertex  $y$  free from color  $c_2$ . Denote  $z$  the other end of edge from  $y$  with color  $c_2$ . If  $z$  is free from color  $c_1$  then we can color  $x, y$  with  $c_2$  and recolor  $y, z$  with  $c_1$ . So we make an alternation. If  $z$  is not free from color,  $c_1$  let's denote  $w$  as the other end of the edge from  $z$  with color  $c_1$ . If  $w$  is free from color  $c_2$ , then again, we can do alternation. We will eventually find an alternating chain because the graph is bipartite. To find the chain, we can use a depth-first search, and each chain contains no more than  $n$  vertices, so the overall complexity is  $O(|E||V|)$ .

Next, we show an easier-to-understand(but slower) algorithm. We first add dummy edges between vertices in  $L$  and  $R$ (maybe with multiplicity) so that every vertex in  $L$  and  $R$  has degree  $\Delta$ . Then we find any perfect matching, color them with the same color, and erase the matching edges, repeating for  $\Delta$  times. It will also yield an edge coloring with  $\Delta$  different colors. The correctness of this algorithm can be proved through Hall's theorem, which we will omit here. The time complexity is  $O(\Delta|E|\sqrt{|V|})$ , which can get accepted if implemented well.

The overall time complexity, translating concerning  $n$ , is  $O(n^3)$  for the first approach and  $O(n^{3.5})$  for the second approach.

## Problem J. Unfair contest

After sorting  $a_1, a_2, \dots, a_{n-1}(b_1, b_2, \dots, b_{n-1})$ , the newly added  $a_n(b_n)$  may fall into three categories, i.e., among the  $s$  highest scores, among the  $t$  lowest scores, or neither. For each of these categories we have a

valid possible range of  $a_n(b_n)$ , the remaining is just careful and thorough case analysis. Be careful with the case  $s = 0$  or  $t = 0$ . The time complexity is  $O(n \log n)$ , due to sorting.

## Problem K. ZYB's kingdom

First, let's consider how to solve the problem assuming  $k = 1$  for each query. The event 1 is equivalent to, for each subtree separated by  $v$ , let  $s$  be the sum of  $c$  in this subtree, then for each vertex  $u$  in the subtree, we add  $s$  to then subtract  $c_u$  from its income.

There exists a folklore  $O(q\sqrt{n} \log n)$  algorithm by considering vertices with degree  $\leq \sqrt{n}$  and  $\geq \sqrt{n}$  separately and using segment tree to update the answer. However, a much more clever approach can solve this in  $O(q \log n)$ . The idea is to first root the tree at any vertex then do a heavy-light decomposition on the tree. Then for each event of type 1, we only update the subtree that is connected to the father of  $v$  and connected to the only heavy child of  $v$ . This takes  $O(\log n)$  time as we only update two subtrees. For each event of type 2, after querying the answer of  $v$  in the segment tree, what are we missing? We need to add the contribution of updates of  $u$  such that  $u$  is an ancestor of  $v$ , but  $v$  is not in the subtree connected to the heavy child of  $u$ . Due to the property of heavy-light decomposition, there are at most  $O(\log n)$  such vertices, so by saving the updates in an array when encountering event of type 1, and iteratively go up along heavy chains when encountering event of type 2, we now can solve the problem in  $O(\log n)$  per event, so the overall complexity is  $O(q \log n)$ .

Then what happens if  $k > 1$ ? Everything is basically the same, you just need to apply the same idea to each component. Here one may need to build a tree structure(a.k.a, virtual tree)with respect to the  $k$  forbidden vertices, then solve recursively on this tree structure. The overall complexity is  $O(\sum k \log n)$ .